

class08

Nate Tran

Data Preparation

```
wisc.df <- read.csv("WisconsinCancer.csv", row.names=1)
wisc.data <- wisc.df[,-1]
diagnosis <- as.factor(wisc.df$diagnosis)
```

Q1

There are 569 observations in the dataset.

```
nrow(wisc.df)
```

```
[1] 569
```

Q2

There are 212 observations with malignant diagnoses.

```
table(diagnosis)
```

```
diagnosis
  B    M
357 212
```

Q3

There are 10 features in the data suffixed with “__mean”.

```
length(grep("_mean", colnames(wisc.df)))
```

```
[1] 10
```

```
#checking if wisc.data needs to be scaled
round(colMeans(wisc.data), 2)
```

| | | |
|------------------------|----------------------|-------------------------|
| radius_mean | texture_mean | perimeter_mean |
| 14.13 | 19.29 | 91.97 |
| area_mean | smoothness_mean | compactness_mean |
| 654.89 | 0.10 | 0.10 |
| concavity_mean | concave.points_mean | symmetry_mean |
| 0.09 | 0.05 | 0.18 |
| fractal_dimension_mean | radius_se | texture_se |
| 0.06 | 0.41 | 1.22 |
| perimeter_se | area_se | smoothness_se |
| 2.87 | 40.34 | 0.01 |
| compactness_se | concavity_se | concave.points_se |
| 0.03 | 0.03 | 0.01 |
| symmetry_se | fractal_dimension_se | radius_worst |
| 0.02 | 0.00 | 16.27 |
| texture_worst | perimeter_worst | area_worst |
| 25.68 | 107.26 | 880.58 |
| smoothness_worst | compactness_worst | concavity_worst |
| 0.13 | 0.25 | 0.27 |
| concave.points_worst | symmetry_worst | fractal_dimension_worst |
| 0.11 | 0.29 | 0.08 |

```
round(apply(wisc.data, 2, sd), 2)
```

| | | |
|------------------------|---------------------|------------------|
| radius_mean | texture_mean | perimeter_mean |
| 3.52 | 4.30 | 24.30 |
| area_mean | smoothness_mean | compactness_mean |
| 351.91 | 0.01 | 0.05 |
| concavity_mean | concave.points_mean | symmetry_mean |
| 0.08 | 0.04 | 0.03 |
| fractal_dimension_mean | radius_se | texture_se |
| 0.01 | 0.28 | 0.55 |
| perimeter_se | area_se | smoothness_se |

| | | | |
|----------------------|----------------------|-------------------------|-------------------|
| | 2.02 | 45.49 | 0.00 |
| compactness_se | | concavity_se | concave.points_se |
| | 0.02 | 0.03 | 0.01 |
| symmetry_se | fractal_dimension_se | | radius_worst |
| | 0.01 | 0.00 | 4.83 |
| texture_worst | perimeter_worst | | area_worst |
| | 6.15 | 33.60 | 569.36 |
| smoothness_worst | compactness_worst | | concavity_worst |
| | 0.02 | 0.16 | 0.21 |
| concave.points_worst | symmetry_worst | fractal_dimension_worst | |
| | 0.07 | 0.06 | 0.02 |

The data needs to be scaled, since the mean and SD values are so variant.

#PCA Time

```
wisc.pr <- prcomp(wisc.data, scale=T)
summary(wisc.pr)
```

Importance of components:

| | PC1 | PC2 | PC3 | PC4 | PC5 | PC6 | PC7 |
|------------------------|---------|---------|---------|---------|---------|---------|---------|
| Standard deviation | 3.6444 | 2.3857 | 1.67867 | 1.40735 | 1.28403 | 1.09880 | 0.82172 |
| Proportion of Variance | 0.4427 | 0.1897 | 0.09393 | 0.06602 | 0.05496 | 0.04025 | 0.02251 |
| Cumulative Proportion | 0.4427 | 0.6324 | 0.72636 | 0.79239 | 0.84734 | 0.88759 | 0.91010 |
| | PC8 | PC9 | PC10 | PC11 | PC12 | PC13 | PC14 |
| Standard deviation | 0.69037 | 0.6457 | 0.59219 | 0.5421 | 0.51104 | 0.49128 | 0.39624 |
| Proportion of Variance | 0.01589 | 0.0139 | 0.01169 | 0.0098 | 0.00871 | 0.00805 | 0.00523 |
| Cumulative Proportion | 0.92598 | 0.9399 | 0.95157 | 0.9614 | 0.97007 | 0.97812 | 0.98335 |
| | PC15 | PC16 | PC17 | PC18 | PC19 | PC20 | PC21 |
| Standard deviation | 0.30681 | 0.28260 | 0.24372 | 0.22939 | 0.22244 | 0.17652 | 0.1731 |
| Proportion of Variance | 0.00314 | 0.00266 | 0.00198 | 0.00175 | 0.00165 | 0.00104 | 0.0010 |
| Cumulative Proportion | 0.98649 | 0.98915 | 0.99113 | 0.99288 | 0.99453 | 0.99557 | 0.9966 |
| | PC22 | PC23 | PC24 | PC25 | PC26 | PC27 | PC28 |
| Standard deviation | 0.16565 | 0.15602 | 0.1344 | 0.12442 | 0.09043 | 0.08307 | 0.03987 |
| Proportion of Variance | 0.00091 | 0.00081 | 0.0006 | 0.00052 | 0.00027 | 0.00023 | 0.00005 |
| Cumulative Proportion | 0.99749 | 0.99830 | 0.9989 | 0.99942 | 0.99969 | 0.99992 | 0.99997 |
| | PC29 | PC30 | | | | | |
| Standard deviation | 0.02736 | 0.01153 | | | | | |
| Proportion of Variance | 0.00002 | 0.00000 | | | | | |
| Cumulative Proportion | 1.00000 | 1.00000 | | | | | |

Q4

44.27% of the original variance is captured by PC1.

Q5

3 PCs are required to describe at least 70% of the original variance.

Q6

7 PCs are required to describe at least 90% of the original variance.

Interpreting PCA

Making scree plot is good way to find “intrinsic dimensionality” of the dataset

```
attributes(wisc.pr)
```

```
$names
```

```
[1] "sdev"      "rotation" "center"    "scale"     "x"
```

```
$class
```

```
[1] "prcomp"
```

```
#calculate variance by squaring SD
```

```
var <- wisc.pr$sdev^2
```


```
#calculate proportion of variance due to each PC
```

```
pve <- var/sum(var)
```


```
plot(pve, xlab = "Principal Component",
```

```
      ylab = "Proportion of Variance Explained",
```

```
      ylim = c(0, 1), type = "o")
```



```
biplot(wisc.pr)
```




Q7

Nothing stands out about this plot aside from its excessive messiness and extreme difficulty in interpreting. This is primarily due to the large amount of observations and variables we are trying to observe here.

Q8

These plots have less and less clear separation between malignant and benign samples when using less important PCs.


```
plot(wisc.pr$x[,c(1,3)], col = diagnosis,  
      xlab = "PC1", ylab = "PC3")
```



Plotting PC plot with ggplot!

```
library(ggplot2)  
  
pc <- as.data.frame(wisc.pr$x)  
  
ggplot(pc) +
```

```
aes(PC1, PC2, col=diagnosis) +  
geom_point()
```



Q9

The feature `concave.points_mean` contributes -0.261 to the first PC.

```
wisc.pr$rotation[,1]
```

| | | |
|------------------------|---------------------|-------------------|
| radius_mean | texture_mean | perimeter_mean |
| -0.21890244 | -0.10372458 | -0.22753729 |
| area_mean | smoothness_mean | compactness_mean |
| -0.22099499 | -0.14258969 | -0.23928535 |
| concavity_mean | concave.points_mean | symmetry_mean |
| -0.25840048 | -0.26085376 | -0.13816696 |
| fractal_dimension_mean | radius_se | texture_se |
| -0.06436335 | -0.20597878 | -0.01742803 |
| perimeter_se | area_se | smoothness_se |
| -0.21132592 | -0.20286964 | -0.01453145 |
| compactness_se | concavity_se | concave.points_se |

| | | |
|----------------------|----------------------|-------------------------|
| -0.17039345 | -0.15358979 | -0.18341740 |
| symmetry_se | fractal_dimension_se | radius_worst |
| -0.04249842 | -0.10256832 | -0.22799663 |
| texture_worst | perimeter_worst | area_worst |
| -0.10446933 | -0.23663968 | -0.22487053 |
| smoothness_worst | compactness_worst | concavity_worst |
| -0.12795256 | -0.21009588 | -0.22876753 |
| concave.points_worst | symmetry_worst | fractal_dimension_worst |
| -0.25088597 | -0.12290456 | -0.13178394 |

Hierarchical clustering

Preparing Data

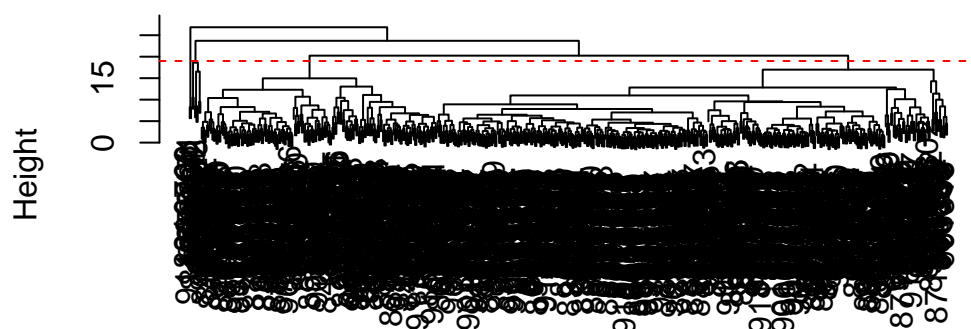
```
data_scaled <- scale(wisc.data)
data_dist <- dist(data_scaled)
wisc.hclust <- hclust(data_dist, method="complete")
```

Q10

With a height of 19, the model using a complete clustering method has 4 clusters.

```
plot(wisc.hclust)
abline(h=19, col="red", lty=2)
```


Cluster Dendrogram



```
data_dist  
hclust (*, "complete")
```

Cutting tree into 4 clusters and assigning membership

```
wisc.clusters <- cutree(wisc.hclust, k=4)  
table(wisc.clusters, diagnosis)
```

| | diagnosis | |
|---------------|-----------|-----|
| wisc.clusters | B | M |
| 1 | 12 | 165 |
| 2 | 2 | 5 |
| 3 | 343 | 40 |
| 4 | 0 | 2 |

Q12


The ward.D2 method gives the best looking results with two more distinct groupings based on the dendrogram, likely correlating with malignant and benign diagnoses.

Combining Methods: PCA and hclust

We can cluster in PC-space using as many or few PCs as we want!

```
#clustering on 3 PC-space
wisc.pr.hclust <- hclust(dist(wisc.pr$x[,1:3]), method="ward.D2")

plot(wisc.pr.hclust)
```



```
dist(wisc.pr$x[, 1:3])
hclust (*, "ward.D2")
```

Finding out whether these two main clusters correlate with malignant and benign

```
groups <- cutree(wisc.pr.hclust, k=2)
table(groups, diagnosis)
```

| | diagnosis | |
|--------|-----------|-----|
| groups | B | M |
| 1 | 24 | 179 |
| 2 | 333 | 33 |

Calculating accuracy of our clustering method with verified diagnoses

```
(179+333)/nrow(wisc.data)
```

```
[1] 0.8998243
```

Q13

The newly created model separates out the two different diagnoses very well with an accuracy of about 90%.

Q14

The other hierarchical clustering models created in previous sections perform similarly to the combined PCA/clustering approach

```
table(wisc.clusters, diagnosis)
```

```
      diagnosis
wisc.clusters  B  M
1      12 165
2       2   5
3     343  40
4       0   2
```

```
(165+343)/nrow(wisc.data)
```

```
[1] 0.8927944
```

#Prediction!

```
new_data <- read.csv("https://tinyurl.com/new-samples-CSV")
new_predict <- predict(wisc.pr, newdata=new_data)
new_predict
```

```
      PC1      PC2      PC3      PC4      PC5      PC6      PC7
[1,]  2.576616 -3.135913  1.3990492 -0.7631950  2.781648 -0.8150185 -0.3959098
[2,] -4.754928 -3.009033 -0.1660946 -0.6052952 -1.140698 -1.2189945  0.8193031
      PC8      PC9      PC10      PC11      PC12      PC13      PC14
[1,] -0.2307350  0.1029569 -0.9272861  0.3411457  0.375921  0.1610764  1.187882
[2,] -0.3307423  0.5281896 -0.4855301  0.7173233 -1.185917  0.5893856  0.303029
      PC15      PC16      PC17      PC18      PC19      PC20
[1,]  0.3216974 -0.1743616 -0.07875393 -0.11207028 -0.08802955 -0.2495216
[2,]  0.1299153  0.1448061 -0.40509706  0.06565549  0.25591230 -0.4289500
      PC21      PC22      PC23      PC24      PC25      PC26
```

```

[1,] 0.1228233 0.09358453 0.08347651 0.1223396 0.02124121 0.078884581
[2,] -0.1224776 0.01732146 0.06316631 -0.2338618 -0.20755948 -0.009833238
      PC27      PC28      PC29      PC30
[1,] 0.220199544 -0.02946023 -0.015620933 0.005269029
[2,] -0.001134152 0.09638361 0.002795349 -0.019015820


```

Plotting new_data predictions against the wisc.pr PC plot

```

plot(wisc.pr$x[,1:2], col= diagnosis)
points(new_predict[,1], new_predict[,2], col="blue", pch=16, cex=3)
text(new_predict[,1], new_predict[,2], c(1,2), col="white")

```



Q16

We should prioritize patient #2 for followup based on the prediction results since patient #2 clusters with other patients with malignant diagnoses.